

```
LINEAR REGRESSION  
DATA
```

In [10]:

```
import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings("ignore")
```

In [8]:

```
data1 = pd.read_csv(r"C:\Users\AKHILA\Downloads\1_2015.csv")
```

In [9]:

```
data1.describe()
```

Out[9]:

	Happiness Rank	Happiness Score	Standard Error	Economy (GDP per Capita)	Family	Health (Life Expectancy)	Freedom
count	158.000000	158.000000	158.000000	158.000000	158.000000	158.000000	158.000000
mean	79.493671	5.375734	0.047885	0.846137	0.991046	0.630259	0.428615
std	45.754363	1.145010	0.017146	0.403121	0.272369	0.247078	0.150693
min	1.000000	2.839000	0.018480	0.000000	0.000000	0.000000	0.000000
25%	40.250000	4.526000	0.037268	0.545808	0.856823	0.439185	0.328330
50%	79.500000	5.232500	0.043940	0.910245	1.029510	0.696705	0.435515
75%	118.750000	6.243750	0.052300	1.158448	1.214405	0.811013	0.549092
max	158.000000	7.587000	0.136930	1.690420	1.402230	1.025250	0.669730



In [11]:

```
data1.head()
```

Out[11]:

	Country	Region	Happiness Rank	Happiness Score	Standard Error	Economy (GDP per Capita)	Family	Health (Life Expectancy)	F
0	Switzerland	Western Europe	1	7.587	0.03411	1.39651	1.34951	0.94143	
1	Iceland	Western Europe	2	7.561	0.04884	1.30232	1.40223	0.94784	
2	Denmark	Western Europe	3	7.527	0.03328	1.32548	1.36058	0.87464	
3	Norway	Western Europe	4	7.522	0.03880	1.45900	1.33095	0.88521	
4	Canada	North America	5	7.427	0.03553	1.32629	1.32261	0.90563	

In [12]:

```
data1.tail()
```

Out[12]:

	Country	Region	Happiness Rank	Happiness Score	Standard Error	Economy (GDP per Capita)	Family	Health (Life Expectancy)	F
153	Rwanda	Sub-Saharan Africa	154	3.465	0.03464	0.22208	0.77370	0.42864	
154	Benin	Sub-Saharan Africa	155	3.340	0.03656	0.28665	0.35386	0.31910	
155	Syria	Middle East and Northern Africa	156	3.006	0.05015	0.66320	0.47489	0.72193	
156	Burundi	Sub-Saharan Africa	157	2.905	0.08658	0.01530	0.41587	0.22396	
157	Togo	Sub-Saharan Africa	158	2.839	0.06727	0.20868	0.13995	0.28443	

In [13]:

```
data1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 158 entries, 0 to 157
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Country                               158 non-null    object
1   Region                                158 non-null    object
2   Happiness Rank                        158 non-null    int64
3   Happiness Score                      158 non-null    float64
4   Standard Error                      158 non-null    float64
5   Economy (GDP per Capita)            158 non-null    float64
6   Family                               158 non-null    float64
7   Health (Life Expectancy)            158 non-null    float64
8   Freedom                              158 non-null    float64
9   Trust (Government Corruption)       158 non-null    float64
10  Generosity                          158 non-null    float64
11  Dystopia Residual                    158 non-null    float64
dtypes: float64(9), int64(1), object(2)
memory usage: 14.9+ KB
```

In [14]:

```
data1.columns
```

Out[14]:

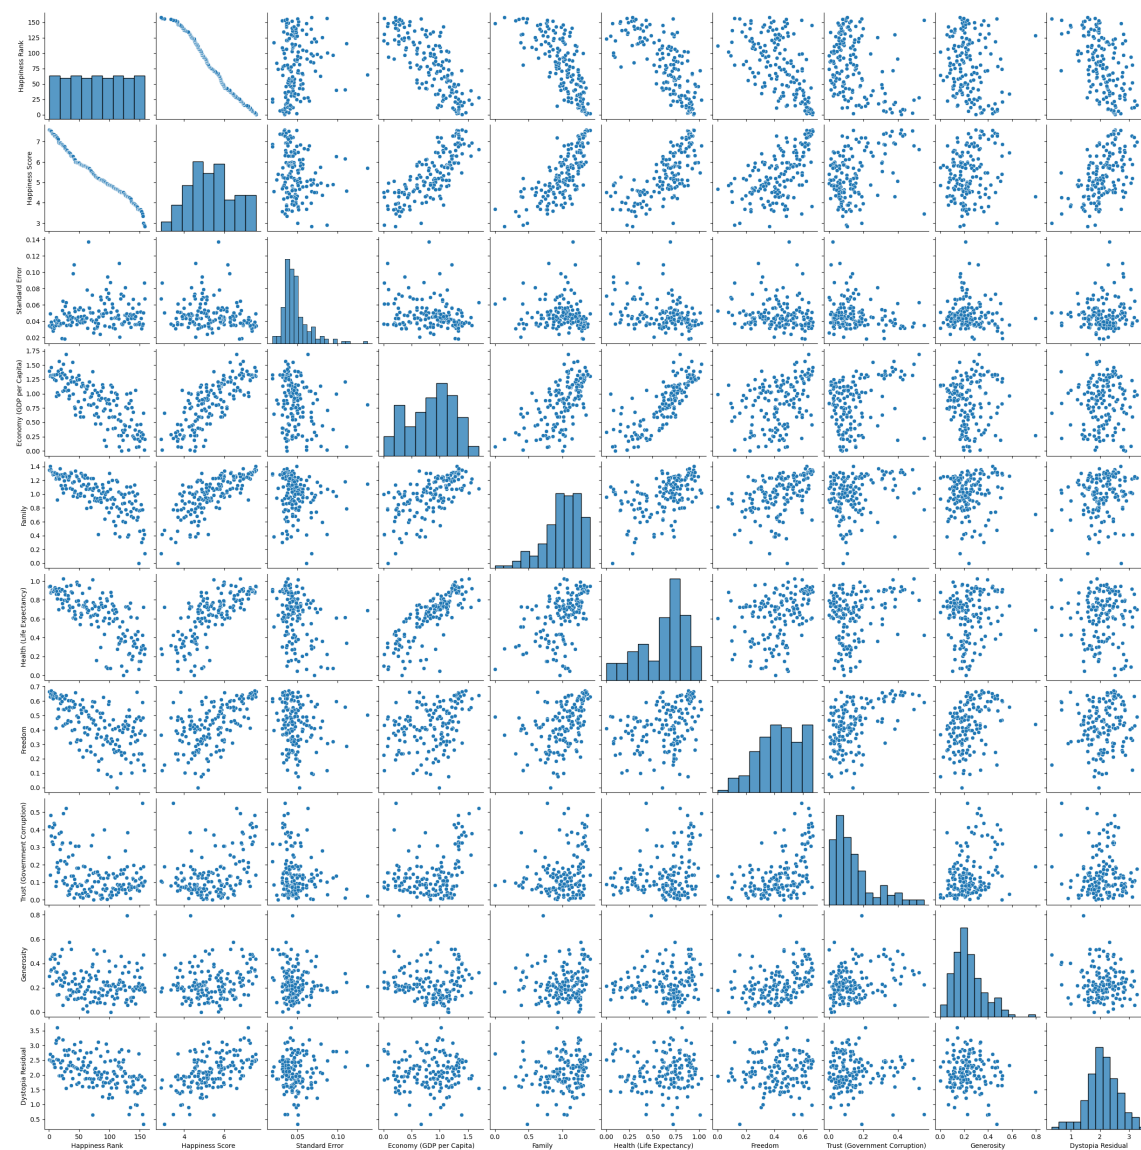
```
Index(['Country', 'Region', 'Happiness Rank', 'Happiness Score',
      'Standard Error', 'Economy (GDP per Capita)', 'Family',
      'Health (Life Expectancy)', 'Freedom', 'Trust (Government Corruptio
n)',
      'Generosity', 'Dystopia Residual'],
      dtype='object')
```

In [15]:

```
sns.pairplot(data1)
```

Out[15]:

<seaborn.axisgrid.PairGrid at 0x1fe4fa25190>




In [16]:

```
sns.displot(data1['Freedom'])
```

Out[16]:

<seaborn.axisgrid.FacetGrid at 0x1fe545c94d0>



In [17]:


```
new_data1 = data1[['Happiness Rank', 'Happiness Score',  
                  'Standard Error', 'Economy (GDP per Capita)', 'Family',  
                  'Health (Life Expectancy)', 'Freedom', 'Trust (Government Corruption)',  
                  'Generosity', 'Dystopia Residual']]
```

In [18]:

```
sns.heatmap(new_data1.corr())
```

Out[18]:

<Axes: >



MODEL BUILDING FOR DATA1

In [19]:

```
X = new_data1[['Happiness Rank', 'Happiness Score',  
               'Standard Error', 'Economy (GDP per Capita)', 'Family',  
               'Health (Life Expectancy)', 'Trust (Government Corruption)',  
               'Generosity', 'Dystopia Residual']]  
y = data1['Freedom']
```

In [21]:

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30)
```

In [22]:

```
from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(X_train,y_train)
```

Out[22]:

```
▼ LinearRegression
LinearRegression()
```

In [24]:

```
#Prediction
predX = lr.predict(X_test)
print(predX)
```

```
[0.62453962 0.48895909 0.3887443  0.66573731 0.41483713 0.25147207
 0.2443236  0.51506286 0.46058424 0.59632617 0.53221435 0.24520173
 0.46858129 0.58419561 0.41685166 0.53489455 0.31865778 0.44011625
 0.53866218 0.39505185 0.53164933 0.42182889 0.38243466 0.46635209
 0.60354277 0.45527026 0.20091261 0.30631689 0.6512724  0.57723779
 0.60322881 0.4229416  0.43440833 0.41626831 0.61755439 0.39724233
 0.26309463 0.5298249  0.49093038 0.64016108 0.10017369 0.17260508
 0.4769477  0.66021869 0.28480765 0.55692497 0.37036783 0.11883106]
```

In [25]:

```
#Accuracy
print(lr.score(X_test,y_test))
```


```
0.9999955087324011
```

In [26]:

```
plt.scatter(y_test,predX)
```

Out[26]:

<matplotlib.collections.PathCollection at 0x1fe5804d190>



DATA 2

In [27]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

In [28]:

```
data2 = pd.read_csv(r"C:\Users\AKHILA\Downloads\4_Drug200.csv")
```



In [29]:

```
data2.describe()
```

Out[29]:

	Age	Na_to_K
count	200.000000	200.000000
mean	44.315000	16.084485
std	16.544315	7.223956
min	15.000000	6.269000
25%	31.000000	10.445500
50%	45.000000	13.936500
75%	58.000000	19.380000
max	74.000000	38.247000

In [30]:

```
data2.head()
```

Out[30]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

In [31]:

```
data2.tail()
```

Out[31]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
195	56	F	LOW	HIGH	11.567	drugC
196	16	M	LOW	HIGH	12.006	drugC
197	52	M	NORMAL	HIGH	9.894	drugX
198	23	M	NORMAL	NORMAL	14.020	drugX
199	40	F	LOW	NORMAL	11.349	drugX

In [32]:

```
data2.columns
```

Out[32]:

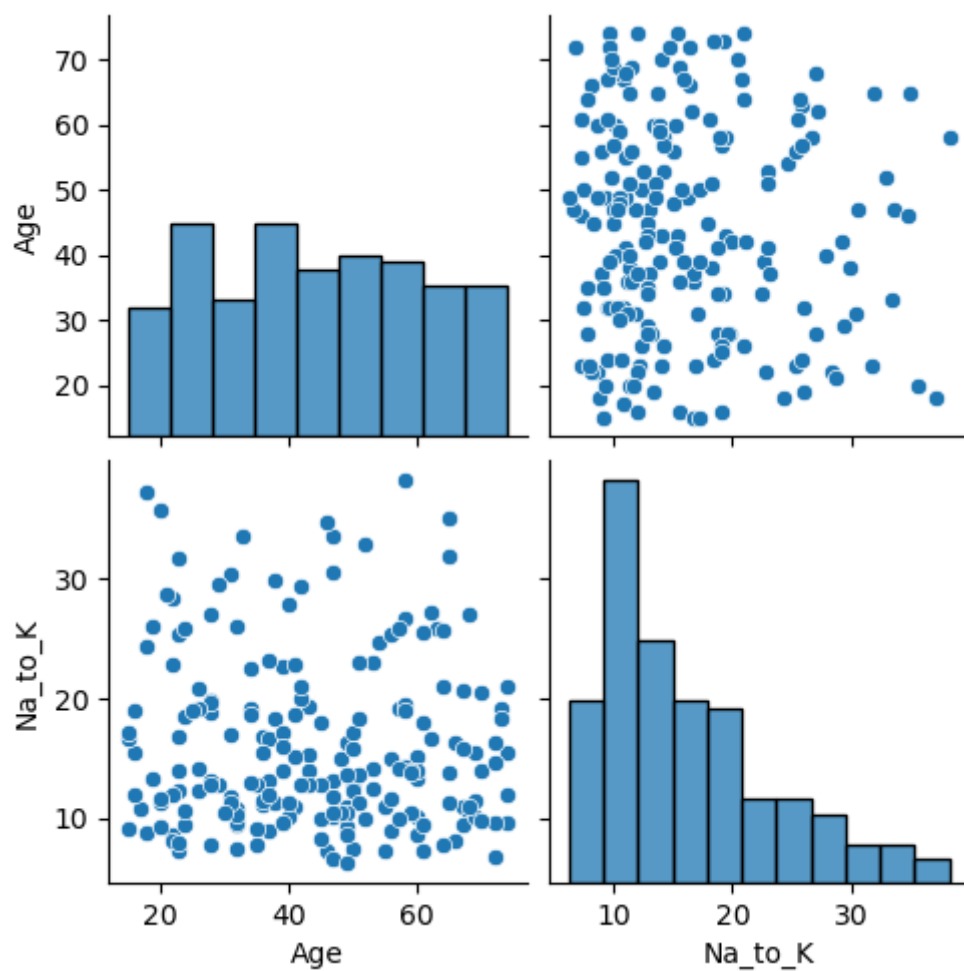
```
Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

In [33]:

```
sns.pairplot(data2)
```

Out[33]:

<seaborn.axisgrid.PairGrid at 0x1fe5801e0d0>



In [34]:

```
#Changing High, Normal and Low to 2, 1 and 0 respectively...
BP = {"BP":{"LOW":0,"NORMAL":1,"HIGH":2}}
data2 = data2.replace(BP)
Cholesterol = {"Cholesterol":{"LOW":0,"NORMAL":1,"HIGH":2}}
data2 = data2.replace(Cholesterol)
data2
```

Out[34]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	2	2	25.355	drugY
1	47	M	0	2	13.093	drugC
2	47	M	0	2	10.114	drugC
3	28	F	1	2	7.798	drugX
4	61	F	0	2	18.043	drugY
...	...	...	...	...	...	...
195	56	F	0	2	11.567	drugC
196	16	M	0	2	12.006	drugC
197	52	M	1	2	9.894	drugX
198	23	M	1	1	14.020	drugX
199	40	F	0	1	11.349	drugX


200 rows × 6 columns

In [35]:

```
sns.displot(data2['BP'])
```

Out[35]:

<seaborn.axisgrid.FacetGrid at 0x1fe580a1650>



In [36]:

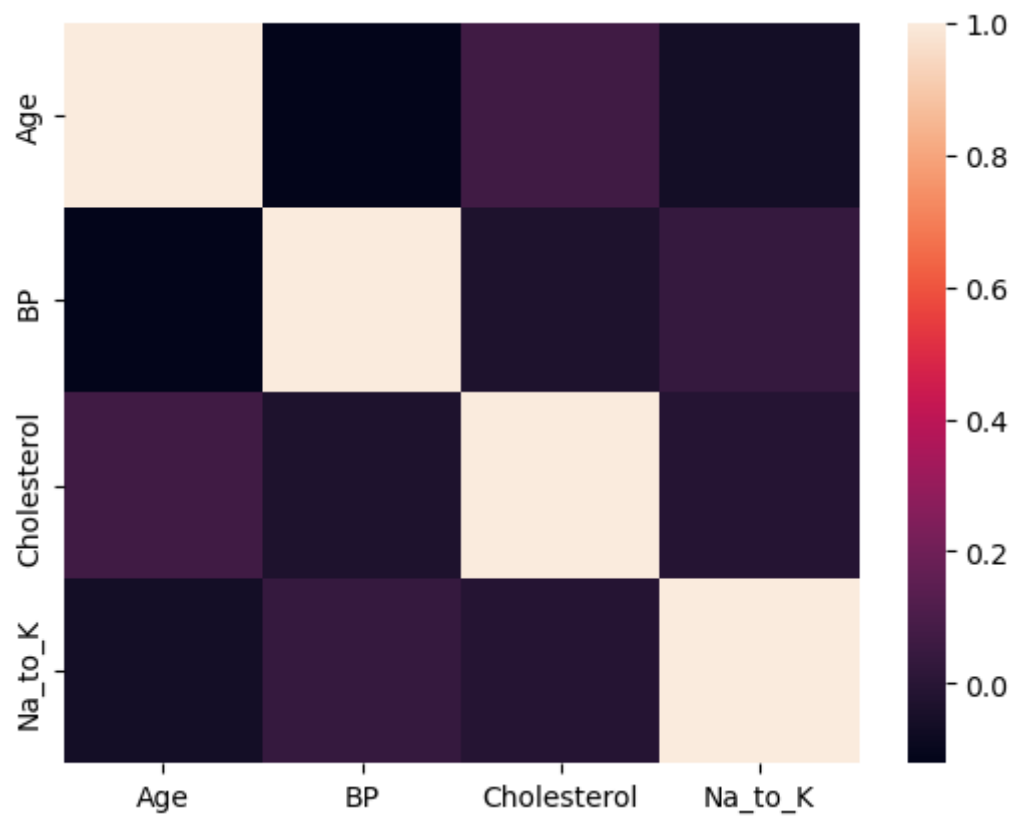
```
new_data2 = data2[['Age', 'BP', 'Cholesterol', 'Na_to_K']]
```

In [37]:

```
sns.heatmap(new_data2.corr())
```

Out[37]:

<Axes: >



MODEL BUILDING FOR DATA2

In [38]:

```
X = new_data2[['Age', 'Cholesterol', 'Na_to_K']]  
y = data2['BP']
```

In [39]:

```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30)
```

In [40]:

```
from sklearn.linear_model import LinearRegression
```

In [41]:

```
lr=LinearRegression()  
lr.fit(X_train,y_train)
```

Out[41]:

```
▼ LinearRegression  
LinearRegression()
```

In [42]:

```
#Prediction  
predX = lr.predict(X_test)  
print(predX)
```

```
[1.23841392 1.23072263 1.03160715 1.23108124 1.22226682 1.05516537  
0.98868878 0.92616347 1.05516375 1.26518257 1.15913418 0.97170118  
0.95071506 1.1920163 0.86384882 0.90074657 1.20160695 1.10109498  
1.24977656 1.11134019 1.1365969 1.14711042 1.00025615 0.91728279  
1.22031298 0.98893481 1.19730414 1.16033698 1.0529568 1.05137644  
1.27602965 1.1605424 1.23058139 1.1892954 1.10247716 1.01907854  
1.11352342 0.91459748 1.1631028 0.92957931 1.25872482 1.12921236  
0.86797911 1.10071873 1.21715257 1.07726067 1.13209191 1.2956641  
1.27438726 0.86892208 1.11292479 1.18576759 1.11521296 1.0193755  
1.1126615 1.18999666 1.31843808 1.01562787 0.96851743 0.96875359]
```

In [43]:

```
#Accuracy  
print(lr.score(X_test,y_test))
```


```
-0.0452325646804046
```

In [44]:

```
plt.scatter(y_test,predX)
```

Out[44]:

<matplotlib.collections.PathCollection at 0x1fe589bb410>



DATA 3

In [48]:

```
data3 = pd.read_csv(r"C:\Users\AKHILA\Downloads\7_Uber.csv")
data3
```

Out[48]:

Unnamed: 0		key	fare_amount	pickup_datetime	pickup_longitude	pick
0	24238194	2015-05-07 19:52:06.0000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	
1	27835199	2009-07-17 20:04:56.0000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	
2	44984355	2009-08-24 21:45:00.00000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	
3	25894730	2009-06-26 08:22:21.0000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	
...	...	...	...	...	...	
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	
199996	16382965	2014-03-14 01:09:00.0000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	
199998	20259894	2015-05-20 14:56:25.0000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	

200000 rows × 9 columns





In [49]:

```
data3.describe()
```

Out[49]:

	Unnamed: 0	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude
count	2.000000e+05	200000.000000	200000.000000	200000.000000	199999.000000	199999.000000
mean	2.771250e+07	11.359955	-72.527638	39.935885	-72.525292	39.935885
std	1.601382e+07	9.901776	11.437787	7.720539	13.117408	7.720539
min	1.000000e+00	-52.000000	-1340.648410	-74.015515	-3356.666300	-74.015515
25%	1.382535e+07	6.000000	-73.992065	40.734796	-73.991407	40.734796
50%	2.774550e+07	8.500000	-73.981823	40.752592	-73.980093	40.752592
75%	4.155530e+07	12.500000	-73.967154	40.767158	-73.963658	40.767158
max	5.542357e+07	499.000000	57.418457	1644.421482	1153.572603	1644.421482

In [50]:

```
data3.head()
```

Out[50]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude
0	24238194	2015-05-07 19:52:06.00000003	7.5	2015-05-07 19:52:06 UTC	-73.999817	40.732451
1	27835199	2009-07-17 20:04:56.00000002	7.7	2009-07-17 20:04:56 UTC	-73.994355	40.732451
2	44984355	2009-08-24 21:45:00.000000061	12.9	2009-08-24 21:45:00 UTC	-74.005043	40.732451
3	25894730	2009-06-26 08:22:21.00000001	5.3	2009-06-26 08:22:21 UTC	-73.976124	40.732451
4	17610152	2014-08-28 17:47:00.000000188	16.0	2014-08-28 17:47:00 UTC	-73.925023	40.732451

In [51]:

```
data3.tail()
```

Out[51]:

	Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	picku
199995	42598914	2012-10-28 10:49:00.00000053	3.0	2012-10-28 10:49:00 UTC	-73.987042	
199996	16382965	2014-03-14 01:09:00.00000008	7.5	2014-03-14 01:09:00 UTC	-73.984722	
199997	27804658	2009-06-29 00:42:00.00000078	30.9	2009-06-29 00:42:00 UTC	-73.986017	
199998	20259894	2015-05-20 14:56:25.00000004	14.5	2015-05-20 14:56:25 UTC	-73.997124	
199999	11951496	2010-05-15 04:08:00.00000076	14.1	2010-05-15 04:08:00 UTC	-73.984395	

In [52]:

```
data3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            200000 non-null int64
1   key                   200000 non-null object
2   fare_amount           200000 non-null float64
3   pickup_datetime       200000 non-null object
4   pickup_longitude      200000 non-null float64
5   pickup_latitude       200000 non-null float64
6   dropoff_longitude     199999 non-null float64
7   dropoff_latitude      199999 non-null float64
8   passenger_count       200000 non-null int64
dtypes: float64(5), int64(2), object(2)
memory usage: 13.7+ MB
```

In [53]:

```
data3.columns
```

Out[53]:

```
Index(['Unnamed: 0', 'key', 'fare_amount', 'pickup_datetime',
      'pickup_longitude', 'pickup_latitude', 'dropoff_longitude',
      'dropoff_latitude', 'passenger_count'],
      dtype='object')
```

```
RANDOM FOREST
```

In [54]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
```

In [55]:

```
df1=pd.read_csv(r"C:\Users\AKHILA\Downloads\C1_Ionosphere.csv")
df1
```

Out[55]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	-
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-
...	...	...	...	...	...	...	...	...	...	...	...	...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	-
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	-
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-

350 rows × 35 columns



In [56]:

```
df1.describe()
```

Out[56]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.0...
count	350.000000	350.0	350.000000	350.000000	350.000000	350.000000	350.000000	350.000000
mean	0.891429	0.0	0.640330	0.044667	0.600350	0.116154	0.549284	0.116154
std	0.311546	0.0	0.498059	0.442032	0.520431	0.461443	0.493124	0.520431
min	0.000000	0.0	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000	-1.000000
25%	1.000000	0.0	0.471517	-0.065388	0.412555	-0.024868	0.209105	-0.024868
50%	1.000000	0.0	0.870795	0.016700	0.808620	0.021170	0.728000	0.021170
75%	1.000000	0.0	1.000000	0.194727	1.000000	0.335317	0.970445	0.335317
max	1.000000	0.0	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

8 rows × 34 columns



In [57]:

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 35 columns):
#   Column                Non-Null Count  Dtype
---  -
0   1                      350 non-null   int64
1   0                      350 non-null   int64
2   0.99539               350 non-null   float64
3   -0.05889              350 non-null   float64
4   0.85243               350 non-null   float64
5   0.02306               350 non-null   float64
6   0.83398               350 non-null   float64
7   -0.37708              350 non-null   float64
8   1.1                   350 non-null   float64
9   0.03760               350 non-null   float64
10  0.85243.1             350 non-null   float64
11  -0.17755              350 non-null   float64
12  0.59755               350 non-null   float64
13  -0.44945              350 non-null   float64
14  0.60536               350 non-null   float64
15  -0.38223              350 non-null   float64
16  0.84356               350 non-null   float64
17  -0.38542              350 non-null   float64
18  0.58212               350 non-null   float64
19  -0.32192              350 non-null   float64
20  0.56971               350 non-null   float64
21  -0.29674              350 non-null   float64
22  0.36946               350 non-null   float64
23  -0.47357              350 non-null   float64
24  0.56811               350 non-null   float64
25  -0.51171              350 non-null   float64
26  0.41078               350 non-null   float64
27  -0.46168              350 non-null   float64
28  0.21266               350 non-null   float64
29  -0.34090              350 non-null   float64
30  0.42267               350 non-null   float64
31  -0.54487              350 non-null   float64
32  0.18641               350 non-null   float64
33  -0.45300              350 non-null   float64
34  g                      350 non-null   object
dtypes: float64(32), int64(2), object(1)
memory usage: 95.8+ KB
```

In [58]:

```
g = {"g":{"g":1,"b":2}}
df1 = df1.replace(g)
df1
```

Out[58]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...	-
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...	-
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...	-
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...	-
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...	-
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...	-
...	...	...	...	...	...	...	...	...	...	...	...	...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...	-
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...	-
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...	-
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...	-
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...	-

350 rows × 35 columns



In [59]:

```
df1["g"].value_counts()
```

Out[59]:

```
1    224
2    126
Name: g, dtype: int64
```

In [60]:

```
x = df1.drop("g",axis=1)
y = df1["g"]
```

In [61]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.40)
```

In [62]:

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[62]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [63]:

```
rf = RandomForestClassifier()
```

In [64]:

```
params = {"max_depth":[1,2,3,4,5],
          "min_samples_leaf":[2,4,6,8,10],
          "n_estimators":[1,3,5,7]
}
```

In [65]:

```
from sklearn.model_selection import GridSearchCV
gs = GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
gs.fit(x_train,y_train)
```

Out[65]:

```
► GridSearchCV
► estimator: RandomForestClassifier
  ► RandomForestClassifier
```

In [66]:

```
rf_best = gs.best_estimator_
rf_best
```

Out[66]:

```
▼ RandomForestClassifier
RandomForestClassifier(max_depth=4, min_samples_leaf=6, n_estimators=7)
```

In [67]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(40,40))
plot_tree(rf_best.estimators_[4],feature_names=None,class_names=['Yes','No'])
```

Out[67]:

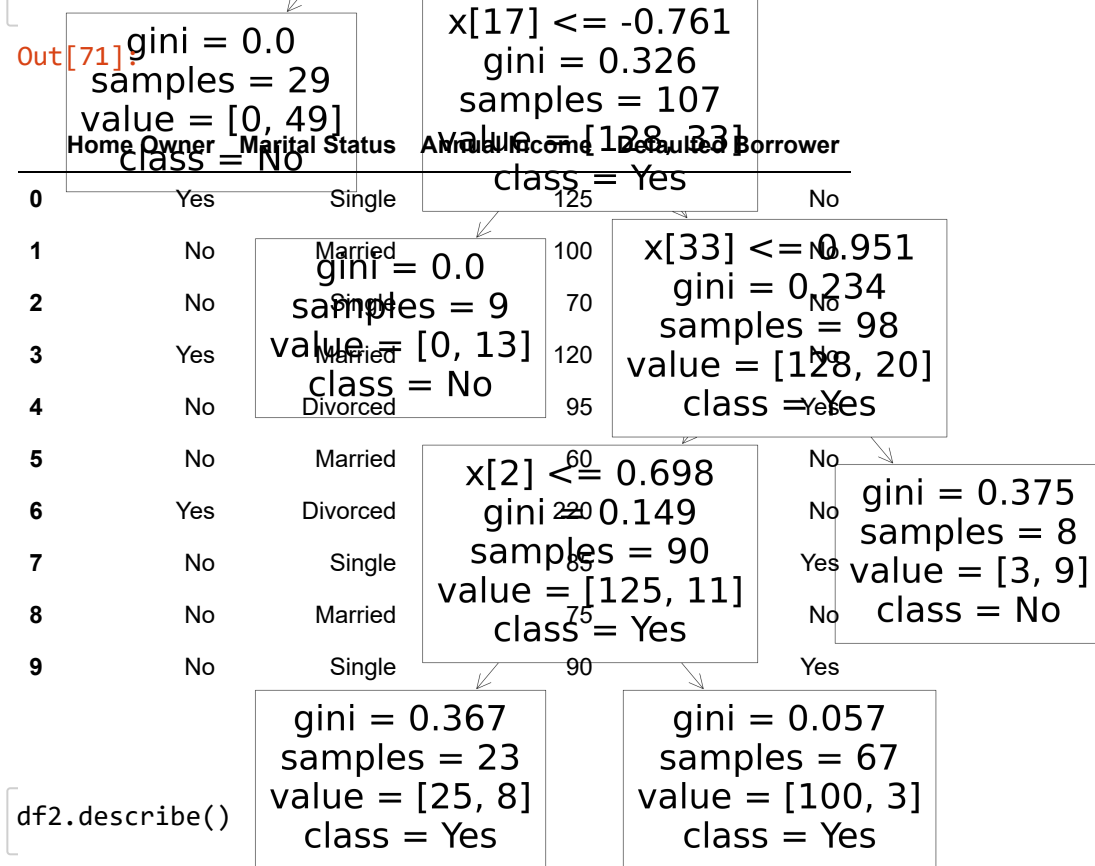
```
[Text(0.3333333333333333, 0.9, 'x[4] <= 0.11\ngini = 0.476\nsamples = 136\nvalue = [128, 82]\nclass = Yes'),
Text(0.16666666666666666, 0.7, 'gini = 0.0\nsamples = 29\nvalue = [0, 49]\nclass = No'),
Text(0.5, 0.7, 'x[17] <= -0.761\ngini = 0.326\nsamples = 107\nvalue = [128, 33]\nclass = Yes'),
Text(0.3333333333333333, 0.5, 'gini = 0.0\nsamples = 9\nvalue = [0, 13]\nclass = No'),
Text(0.6666666666666666, 0.5, 'x[33] <= 0.951\ngini = 0.234\nsamples = 98\nvalue = [128, 20]\nclass = Yes'),
Text(0.5, 0.3, 'x[2] <= 0.698\ngini = 0.149\nsamples = 90\nvalue = [125, 11]\nclass = Yes'),
Text(0.3333333333333333, 0.1, 'gini = 0.367\nsamples = 23\nvalue = [25, 8]\nclass = Yes'),
Text(0.6666666666666666, 0.1, 'gini = 0.057\nsamples = 67\nvalue = [100, 3]\nclass = Yes'),
Text(0.8333333333333334, 0.3, 'gini = 0.375\nsamples = 8\nvalue = [3, 9]\nclass = No')]
```



```

RANDOM FOREST FROM DATA 2
x[4] <= 0.11
gini = 0.476
samples = 136
value = [128, 82]
class = Yes
df2=pd.read_csv(r"C:\Users\AKSLA\Downloads\C10_Loan1.csv")
df2

```



Out[72]:

Annual Income	
count	10.000000
mean	104.000000
std	45.631373
min	60.000000
25%	77.500000
50%	92.500000
75%	115.000000
max	220.000000

In [73]:

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10 entries, 0 to 9
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Home Owner             10 non-null    object
1   Marital Status         10 non-null    object
2   Annual Income          10 non-null    int64
3   Defaulted Borrower     10 non-null    object
dtypes: int64(1), object(3)
memory usage: 452.0+ bytes
```

In [74]:

```
Home_Owner = {"Home Owner":{"Yes":1,"No":2}}
df2 = df2.replace(Home_Owner)
Defaulted_Borrower = {"Defaulted Borrower":{"Yes":1,"No":2}}
df2 = df2.replace(Defaulted_Borrower)
Marital_Status = {"Marital Status":{"Divorced":0,"Single":1,"Married":2}}
df2 = df2.replace(Marital_Status)
df2
```

Out[74]:

	Home Owner	Marital Status	Annual Income	Defaulted Borrower
0	1	1	125	2
1	2	2	100	2
2	2	1	70	2
3	1	2	120	2
4	2	0	95	1
5	2	2	60	2
6	1	0	220	2
7	2	1	85	1
8	2	2	75	2
9	2	1	90	1

In [75]:

```
df2["Home Owner"].value_counts()
```

Out[75]:

```
2    7
1    3
Name: Home Owner, dtype: int64
```

In [76]:

```
df2["Defaulted Borrower"].value_counts()
```

Out[76]:

```
2    7
1    3
Name: Defaulted Borrower, dtype: int64
```

In [77]:

```
x = df2.drop("Marital Status",axis=1)
y = df2["Marital Status"]
```

In [78]:

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.40)
```

In [79]:

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[79]:

```
▼ RandomForestClassifier
RandomForestClassifier()
```

In [80]:

```
rf = RandomForestClassifier()
```

In [81]:

```
params = {"max_depth":[1,2,3,4,5],
          "min_samples_leaf":[2,4,6,8,10],
          "n_estimators":[1,3,5,7]
}
```

In [82]:

```
from sklearn.model_selection import GridSearchCV
gs = GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring='accuracy')
gs.fit(x_train,y_train)
```

Out[82]:

```
► GridSearchCV
► estimator: RandomForestClassifier
  ► RandomForestClassifier
```

In [83]:

```
rf_best = gs.best_estimator_  
rf_best
```

Out[83]:

```
RandomForestClassifier  
RandomForestClassifier(max_depth=1, min_samples_leaf=2, n_estimators=5)
```

In [84]:

```
from sklearn.tree import plot_tree  
plt.figure(figsize=(40,40))  
plot_tree(rf_best.estimators_[4], feature_names=None, class_names=['Yes', 'No'])
```

Out[84]:

```
[Text(0.5, 0.5, 'gini = 0.444\nsamples = 4\nvalue = [0, 4, 2]\n      class = No')]
```

gini = 0.444  
samples = 4  
value = [0, 4, 2]  
class = No

In [ ]: